

Linear Regression on Automobiles

```
In [ ]: !pip install ucimlrepo
```

```
Collecting ucimlrepo  
  Downloading ucimlrepo-0.0.6-py3-none-any.whl (8.0 kB)  
Installing collected packages: ucimlrepo  
Successfully installed ucimlrepo-0.0.6
```

```
In [ ]: #importing libraries  
import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression  
from sklearn.metrics import mean_squared_error, r2_score  
from sklearn.preprocessing import StandardScaler  
from scipy import stats
```

```
In [ ]: from ucimlrepo import fetch_ucirepo  
  
# fetch dataset  
automobile = fetch_ucirepo(id=10)  
  
# data (as pandas dataframes)  
X = automobile.data.features  
y = automobile.data.targets  
  
# metadata  
print(automobile.metadata)  
  
# variable information  
print(automobile.variables)
```

```
{'uci_id': 10, 'name': 'Automobile', 'repository_url': 'https://archive.ics.uci.edu/
dataset/10/automobile', 'data_url': 'https://archive.ics.uci.edu/static/public/10/da
ta.csv', 'abstract': "From 1985 Ward's Automotive Yearbook", 'area': 'Other', 'task
s': ['Regression'], 'characteristics': ['Multivariate'], 'num_instances': 205, 'num_
features': 25, 'feature_types': ['Categorical', 'Integer', 'Real'], 'demographics':
[], 'target_col': ['symboling'], 'index_col': None, 'has_missing_values': 'yes', 'mi
ssing_values_symbol': 'NaN', 'year_of_dataset_creation': 1985, 'last_updated': 'Thu
Aug 10 2023', 'dataset_doi': '10.24432/C5B01C', 'creators': ['Jeffrey Schlimmer'],
'intro_paper': None, 'additional_info': {'summary': 'This data set consists of three
types of entities: (a) the specification of an auto in terms of various characterist
ics, (b) its assigned insurance risk rating, (c) its normalized losses in use as com
pared to other cars. The second rating corresponds to the degree to which the auto
is more risky than its price indicates. Cars are initially assigned a risk factor sy
mbol associated with its price. Then, if it is more risky (or less), this symbol i
s adjusted by moving it up (or down) the scale. Actuarians call this process "symbo
ling". A value of +3 indicates that the auto is risky, -3 that it is probably prett
y safe.\r\n\r\nThe third factor is the relative average loss payment per insured veh
icle year. This value is normalized for all autos within a particular size classifi
cation (two-door small, station wagons, sports/speciality, etc...), and represents t
he average loss per car per year.\r\n\r\nNote: Several of the attributes in the data
base could be used as a "class" attribute.', 'purpose': None, 'funded_by': None, 'in
stances_represent': None, 'recommended_data_splits': None, 'sensitive_data': None,
'preprocessing_description': None, 'variable_info': 'Attribute: Attribute Range\r\n
\r\n 1. symboling: -3, -2, -1, 0, 1, 2, 3.\r\n 2. normalized-losse
s: continuous from 65 to 256.\r\n 3. make: \r\n
alfa-romero, audi, bmw, chevrolet, dodge, honda,\r\n i
suzu, jaguar, mazda, mercedes-benz, mercury,\r\n mitsu
bishi, nissan, peugot, plymouth, porsche,\r\n renault,
saab, subaru, toyota, volkswagen, volvo\r\n\r\n 4. fuel-type: diese
l, gas.\r\n 5. aspiration: std, turbo.\r\n 6. num-of-doors:
four, two.\r\n 7. body-style: hardtop, wagon, sedan, hatchback, conve
rtible.\r\n 8. drive-wheels: 4wd, fwd, rwd.\r\n 9. engine-location:
front, rear.\r\n 10. wheel-base: continuous from 86.6 120.9.\r\n 11. l
ength: continuous from 141.1 to 208.1.\r\n 12. width:
continuous from 60.3 to 72.3.\r\n 13. height: continuous from 47.8
to 59.8.\r\n 14. curb-weight: continuous from 1488 to 4066.\r\n 15. eng
ine-type: dohc, dohcvt, l, ohc, ohcvt, ohcvt, rotor.\r\n 16. num-of-cylind
ers: eight, five, four, six, three, twelve, two.\r\n 17. engine-size:
continuous from 61 to 326.\r\n 18. fuel-system: 1bbl, 2bbl, 4bbl, idi,
mfi, mpfi, spdi, spfi.\r\n 19. bore: continuous from 2.54 to 3.9
4.\r\n 20. stroke: continuous from 2.07 to 4.17.\r\n 21. compressi
on-ratio: continuous from 7 to 23.\r\n 22. horsepower: continuo
us from 48 to 288.\r\n 23. peak-rpm: continuous from 4150 to 6600.\r
\r\n 24. city-mpg: continuous from 13 to 49.\r\n 25. highway-mpg:
continuous from 16 to 54.\r\n 26. price: continuous from 5118 to
45400.', 'citation': None}}
```

	name	role	type	demographic	\
0	price	Feature	Continuous	None	
1	highway-mpg	Feature	Continuous	None	
2	city-mpg	Feature	Continuous	None	
3	peak-rpm	Feature	Continuous	None	
4	horsepower	Feature	Continuous	None	
5	compression-ratio	Feature	Continuous	None	
6	stroke	Feature	Continuous	None	
7	bore	Feature	Continuous	None	
8	fuel-system	Feature	Categorical	None	

9	engine-size	Feature	Continuous	None
10	num-of-cylinders	Feature	Integer	None
11	engine-type	Feature	Categorical	None
12	curb-weight	Feature	Continuous	None
13	height	Feature	Continuous	None
14	width	Feature	Continuous	None
15	length	Feature	Continuous	None
16	wheel-base	Feature	Continuous	None
17	engine-location	Feature	Binary	None
18	drive-wheels	Feature	Categorical	None
19	body-style	Feature	Categorical	None
20	num-of-doors	Feature	Integer	None
21	aspiration	Feature	Binary	None
22	fuel-type	Feature	Binary	None
23	make	Feature	Categorical	None
24	normalized-losses	Feature	Continuous	None
25	symboling	Target	Integer	None

		description	units	missing_values
0		continuous from 5118 to 45400	None	yes
1		continuous from 16 to 54	None	no
2		continuous from 13 to 49	None	no
3		continuous from 4150 to 6600	None	yes
4		continuous from 48 to 288	None	yes
5		continuous from 7 to 23	None	no
6		continuous from 2.07 to 4.17	None	yes
7		continuous from 2.54 to 3.94	None	yes
8	1bbl, 2bbl, 4bbl, idi, mfi, mpfi, spdi, spfi		None	no
9		continuous from 61 to 326	None	no
10	eight, five, four, six, three, twelve, two		None	no
11	dohc, dohcv, l, ohc, ohcf, ohcv, rotor		None	no
12		continuous from 1488 to 4066	None	no
13		continuous from 47.8 to 59.8	None	no
14		continuous from 60.3 to 72.3	None	no
15		continuous from 141.1 to 208.1	None	no
16		continuous from 86.6 120.9	None	no
17		front, rear	None	no
18		4wd, fwd, rwd	None	no
19	hardtop, wagon, sedan, hatchback, convertible		None	no
20		four, two	None	yes
21		std, turbo	None	no
22		diesel, gas	None	no
23	alfa-romero, audi, bmw, chevrolet, dodge, hond...		None	no
24		continuous from 65 to 256	None	yes
25		-3, -2, -1, 0, 1, 2, 3	None	no

```
In [ ]: X.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 201 non-null   float64
1   highway-mpg          205 non-null   int64
2   city-mpg             205 non-null   int64
3   peak-rpm             203 non-null   float64
4   horsepower            203 non-null   float64
5   compression-ratio    205 non-null   float64
6   stroke               201 non-null   float64
7   bore                 201 non-null   float64
8   fuel-system          205 non-null   object
9   engine-size          205 non-null   int64
10  num-of-cylinders     205 non-null   int64
11  engine-type          205 non-null   object
12  curb-weight          205 non-null   int64
13  height               205 non-null   float64
14  width                205 non-null   float64
15  length               205 non-null   float64
16  wheel-base           205 non-null   float64
17  engine-location      205 non-null   object
18  drive-wheels         205 non-null   object
19  body-style           205 non-null   object
20  num-of-doors         203 non-null   float64
21  aspiration           205 non-null   object
22  fuel-type            205 non-null   object
23  make                 205 non-null   object
24  normalized-losses    164 non-null   float64
dtypes: float64(12), int64(5), object(8)
memory usage: 40.2+ KB

```

```
In [ ]: print(automobile.variables)
```

	name	role	type	demographic \
0	price	Feature	Continuous	None
1	highway-mpg	Feature	Continuous	None
2	city-mpg	Feature	Continuous	None
3	peak-rpm	Feature	Continuous	None
4	horsepower	Feature	Continuous	None
5	compression-ratio	Feature	Continuous	None
6	stroke	Feature	Continuous	None
7	bore	Feature	Continuous	None
8	fuel-system	Feature	Categorical	None
9	engine-size	Feature	Continuous	None
10	num-of-cylinders	Feature	Integer	None
11	engine-type	Feature	Categorical	None
12	curb-weight	Feature	Continuous	None
13	height	Feature	Continuous	None
14	width	Feature	Continuous	None
15	length	Feature	Continuous	None
16	wheel-base	Feature	Continuous	None
17	engine-location	Feature	Binary	None
18	drive-wheels	Feature	Categorical	None
19	body-style	Feature	Categorical	None
20	num-of-doors	Feature	Integer	None
21	aspiration	Feature	Binary	None
22	fuel-type	Feature	Binary	None
23	make	Feature	Categorical	None
24	normalized-losses	Feature	Continuous	None
25	symboling	Target	Integer	None

		description	units	missing_values
0		continuous from 5118 to 45400	None	yes
1		continuous from 16 to 54	None	no
2		continuous from 13 to 49	None	no
3		continuous from 4150 to 6600	None	yes
4		continuous from 48 to 288	None	yes
5		continuous from 7 to 23	None	no
6		continuous from 2.07 to 4.17	None	yes
7		continuous from 2.54 to 3.94	None	yes
8	1bbl, 2bbl, 4bbl, idi, mfi, mpfi, spdi, spfi		None	no
9		continuous from 61 to 326	None	no
10	eight, five, four, six, three, twelve, two		None	no
11	dohc, dohcv, l, ohc, ohcf, ohcv, rotor		None	no
12		continuous from 1488 to 4066	None	no
13		continuous from 47.8 to 59.8	None	no
14		continuous from 60.3 to 72.3	None	no
15		continuous from 141.1 to 208.1	None	no
16		continuous from 86.6 120.9	None	no
17		front, rear	None	no
18		4wd, fwd, rwd	None	no
19	hardtop, wagon, sedan, hatchback, convertible		None	no
20		four, two	None	yes
21		std, turbo	None	no
22		diesel, gas	None	no
23	alfa-romero, audi, bmw, chevrolet, dodge, hond...		None	no
24		continuous from 65 to 256	None	yes
25		-3, -2, -1, 0, 1, 2, 3	None	no

```
In [ ]: print(automobile.data.columns)
```

None

```
In [ ]: from ucimlrepo import fetch_ucirepo

# Fetch dataset
automobile = fetch_ucirepo(id=10)

if automobile is None:
    print("Error: Dataset loading failed.")
else:
    print("Dataset loaded successfully.")
```

Dataset loaded successfully.

Loading the dataset

```
In [ ]: automobile = fetch_ucirepo(id=10)
X = automobile.data.features
y = automobile.data.target
```

Data Wrangling

```
In [ ]: #checks for missing values
print(new_X.isnull().sum())
```

```
price                4
highway-mpg          0
city-mpg             0
peak-rpm             2
horsepower           2
..
make_toyota          0
make_volkswagen      0
make_volvo           0
engine-location_front 0
engine-location_rear  0
Length: 68, dtype: int64
```

```
In [ ]: #encodes categorical variables
new_X = pd.get_dummies(X, columns=['fuel-system', 'engine-type', 'drive-wheels', 'b

print(new_X.shape)
print(new_X.dtypes)
```

```
(205, 68)
price                float64
highway-mpg          int64
city-mpg              int64
peak-rpm              float64
horsepower            float64
...
make_toyota           bool
make_volkswagen        bool
make_volvo             bool
engine-location_front  bool
engine-location_rear   bool
Length: 68, dtype: object
```

Feature Scaling:

```
In [ ]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

numeric_columns = new_X.select_dtypes(include=['float64', 'int64']).columns
X_scaled = new_X.copy()
X_scaled[numeric_columns] = scaler.fit_transform(new_X[numeric_columns])
```

Data Exploration

```
In [ ]: X.describe()
```

```
Out[ ]:
```

	price	highway-mpg	city-mpg	peak-rpm	horsepower	compression-ratio	
count	201.000000	205.000000	205.000000	203.000000	203.000000	205.000000	201.0
mean	13207.129353	30.751220	25.219512	5125.369458	104.256158	10.142537	3.0
std	7947.066342	6.886443	6.542142	479.334560	39.714369	3.972040	0.0
min	5118.000000	16.000000	13.000000	4150.000000	48.000000	7.000000	2.0
25%	7775.000000	25.000000	19.000000	4800.000000	70.000000	8.600000	3.0
50%	10295.000000	30.000000	24.000000	5200.000000	95.000000	9.000000	3.0
75%	16500.000000	34.000000	30.000000	5500.000000	116.000000	9.400000	3.0
max	45400.000000	54.000000	49.000000	6600.000000	288.000000	23.000000	4.0

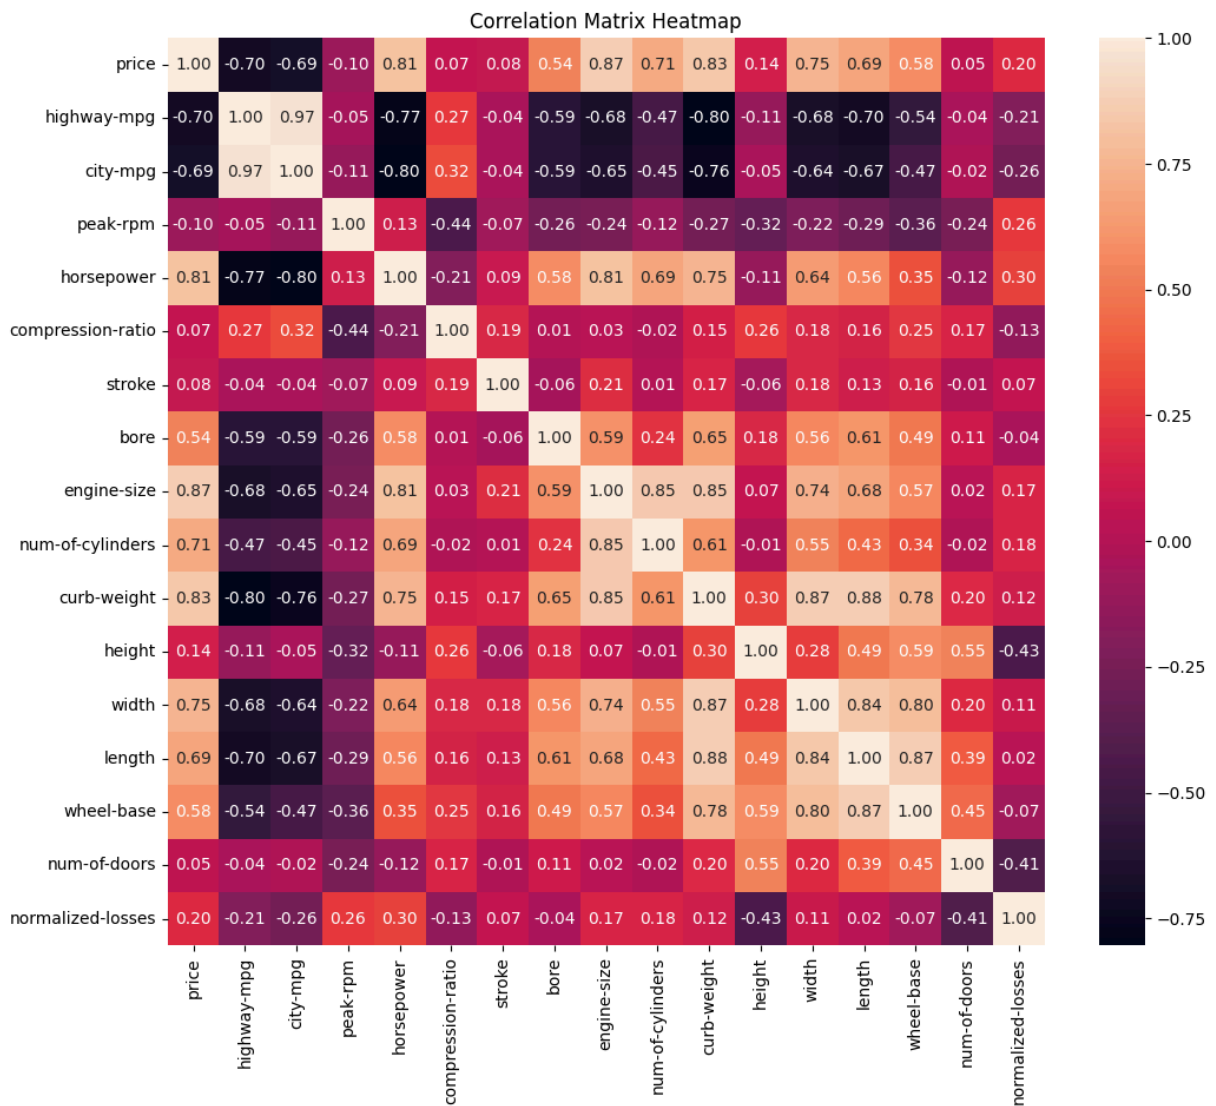
```
In [ ]: #correlatiioin analysis
print(new_X.dtypes)

data_numeric = new_X.select_dtypes(include=['float64', 'int64'])

correlation_matrix = data_numeric.corr()
```

```
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='rocket', fmt=".2f")
plt.title("Correlation Matrix Heatmap")
plt.show()
```

```
price                float64
highway-mpg          int64
city-mpg              int64
peak-rpm              float64
horsepower            float64
...
make_toyota           bool
make_volkswagen        bool
make_volvo             bool
engine-location_front  bool
engine-location_rear   bool
Length: 68, dtype: object
```



Simple Linear Regression

Split Data into Training and Testing Sets:

```
In [ ]: #splitting the data into training and testing sets
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, ra
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-17-5ebc7f651068> in <cell line: 3>()
      1 from sklearn.model_selection import train_test_split
      2
----> 3 X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=
0.2, random_state=42)

NameError: name 'X_encoded' is not defined
```

```
In [ ]: from sklearn.linear_model import LinearRegression

model = LinearRegression()

model.fit(X_train, y_train)
```

```
Out[ ]: ▾ LinearRegression
LinearRegression()
```

Fit Linear Regression Model:

```
In [ ]: from sklearn.metrics import mean_squared_error, r2_score

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")
```

Mean Squared Error: 0.6859129086119733
R-squared: 0.5320537340191854

Logistic Regression on Wine

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [ ]: from ucimlrepo import fetch_ucirepo

# fetch dataset
wine = fetch_ucirepo(id=109)

# data (as pandas dataframes)
X = wine.data.features
y = wine.data.targets

# metadata
print(wine.metadata)

# variable information
print(wine.variables)
```

```
{'uci_id': 109, 'name': 'Wine', 'repository_url': 'https://archive.ics.uci.edu/dataset/109/wine', 'data_url': 'https://archive.ics.uci.edu/static/public/109/data.csv', 'abstract': 'Using chemical analysis to determine the origin of wines', 'area': 'Physics and Chemistry', 'tasks': ['Classification'], 'characteristics': ['Tabular'], 'num_instances': 178, 'num_features': 13, 'feature_types': ['Integer', 'Real'], 'demographics': [], 'target_col': ['class'], 'index_col': None, 'has_missing_values': 'no', 'missing_values_symbol': None, 'year_of_dataset_creation': 1992, 'last_updated': 'Mon Aug 28 2023', 'dataset_doi': '10.24432/C5PC7J', 'creators': ['Stefan Aeberhard', 'M. Forina'], 'intro_paper': {'title': 'Comparative analysis of statistical pattern recognition methods in high dimensional settings', 'authors': 'S. Aeberhard, D. Coomans, O. Vel', 'published_in': 'Pattern Recognition', 'year': 1994, 'url': 'https://www.semanticscholar.org/paper/83dc3e4030d7b9fbdbb4bde03ce12ab70ca10528', 'doi': '10.1016/0031-3203(94)90145-7'}, 'additional_info': {'summary': 'These data are the results of a chemical analysis of wines grown in the same region in Italy but derived from three different cultivars. The analysis determined the quantities of 13 constituents found in each of the three types of wines. \r\n\r\nI think that the initial data set had around 30 variables, but for some reason I only have the 13 dimensional version. I had a list of what the 30 or so variables were, but a.) I lost it, and b.), I would not know which 13 variables are included in the set.\r\n\r\nThe attributes are (donated by Riccardo Leardi, riclea@anchem.unige.it) \r\n1) Alcohol\r\n2) Malic acid\r\n3) Ash\r\n4) Alkalinity of ash \r\n5) Magnesium\r\n6) Total phenols\r\n7) Flavanoids\r\n8) Nonflavanoid phenols\r\n9) Proanthocyanins\r\n10) Color intensity\r\n11) Hue\r\n12) OD280/OD315 of diluted wines\r\n13) Proline \r\n\r\nIn a classification context, this is a well posed problem with "well behaved" class structures. A good data set for first testing of a new classifier, but not very challenging.', 'purpose': 'test', 'funded_by': None, 'instances_represent': None, 'recommended_data_splits': None, 'sensitive_data': None, 'preprocessing_description': None, 'variable_info': 'All attributes are continuous\r\n\r\nNo statistics available, but suggest to standardise variables for certain uses (e.g. for us with classifiers which are NOT scale invariant)\r\n\r\nNOTE: 1st attribute is class identifier (1-3)', 'citation': None}}
```

	name	role	type	demographic	\
0	class	Target	Categorical	None	
1	Alcohol	Feature	Continuous	None	
2	Malicacid	Feature	Continuous	None	
3	Ash	Feature	Continuous	None	
4	Alkalinity_of_ash	Feature	Continuous	None	
5	Magnesium	Feature	Integer	None	
6	Total_phenols	Feature	Continuous	None	
7	Flavanoids	Feature	Continuous	None	
8	Nonflavanoid_phenols	Feature	Continuous	None	
9	Proanthocyanins	Feature	Continuous	None	
10	Color_intensity	Feature	Continuous	None	
11	Hue	Feature	Continuous	None	
12	OD280_OD315_of_diluted_wines	Feature	Continuous	None	
13	Proline	Feature	Integer	None	

	description	units	missing_values
0	None	None	no
1	None	None	no
2	None	None	no
3	None	None	no
4	None	None	no
5	None	None	no
6	None	None	no
7	None	None	no

8	None	None	no
9	None	None	no
10	None	None	no
11	None	None	no
12	None	None	no
13	None	None	no

Loading the dataset

```
In [ ]: wine = pd.DataFrame(data=X, columns=wine.variables['name'][1:])  
wine['class'] = y
```

Data Exploration

```
In [ ]: print(wine.head())  
print(wine.describe())  
print(wine.info())
```

name	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	\
0	14.23	1.71	2.43	15.6	127	2.80	
1	13.20	1.78	2.14	11.2	100	2.65	
2	13.16	2.36	2.67	18.6	101	2.80	
3	14.37	1.95	2.50	16.8	113	3.85	
4	13.24	2.59	2.87	21.0	118	2.80	

name	Flavanoids	Nonflavanoid_phenols	Proanthocyanins	Color_intensity	\
0	3.06		0.28	2.29	5.64
1	2.76		0.26	1.28	4.38
2	3.24		0.30	2.81	5.68
3	3.49		0.24	2.18	7.80
4	2.69		0.39	1.82	4.32

name	Hue	0D280_0D315_of_diluted_wines	Proline	class
0	1.04	3.92	1065	1
1	1.05	3.40	1050	1
2	1.03	3.17	1185	1
3	0.86	3.45	1480	1
4	1.04	2.93	735	1

name	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	\
count	178.000000	178.000000	178.000000	178.000000	178.000000	
mean	13.000618	2.336348	2.366517	19.494944	99.741573	
std	0.811827	1.117146	0.274344	3.339564	14.282484	
min	11.030000	0.740000	1.360000	10.600000	70.000000	
25%	12.362500	1.602500	2.210000	17.200000	88.000000	
50%	13.050000	1.865000	2.360000	19.500000	98.000000	
75%	13.677500	3.082500	2.557500	21.500000	107.000000	
max	14.830000	5.800000	3.230000	30.000000	162.000000	

name	Total_phenols	Flavanoids	Nonflavanoid_phenols	Proanthocyanins	\
count	178.000000	178.000000	178.000000	178.000000	
mean	2.295112	2.029270	0.361854	1.590899	
std	0.625851	0.998859	0.124453	0.572359	
min	0.980000	0.340000	0.130000	0.410000	
25%	1.742500	1.205000	0.270000	1.250000	
50%	2.355000	2.135000	0.340000	1.555000	
75%	2.800000	2.875000	0.437500	1.950000	
max	3.880000	5.080000	0.660000	3.580000	

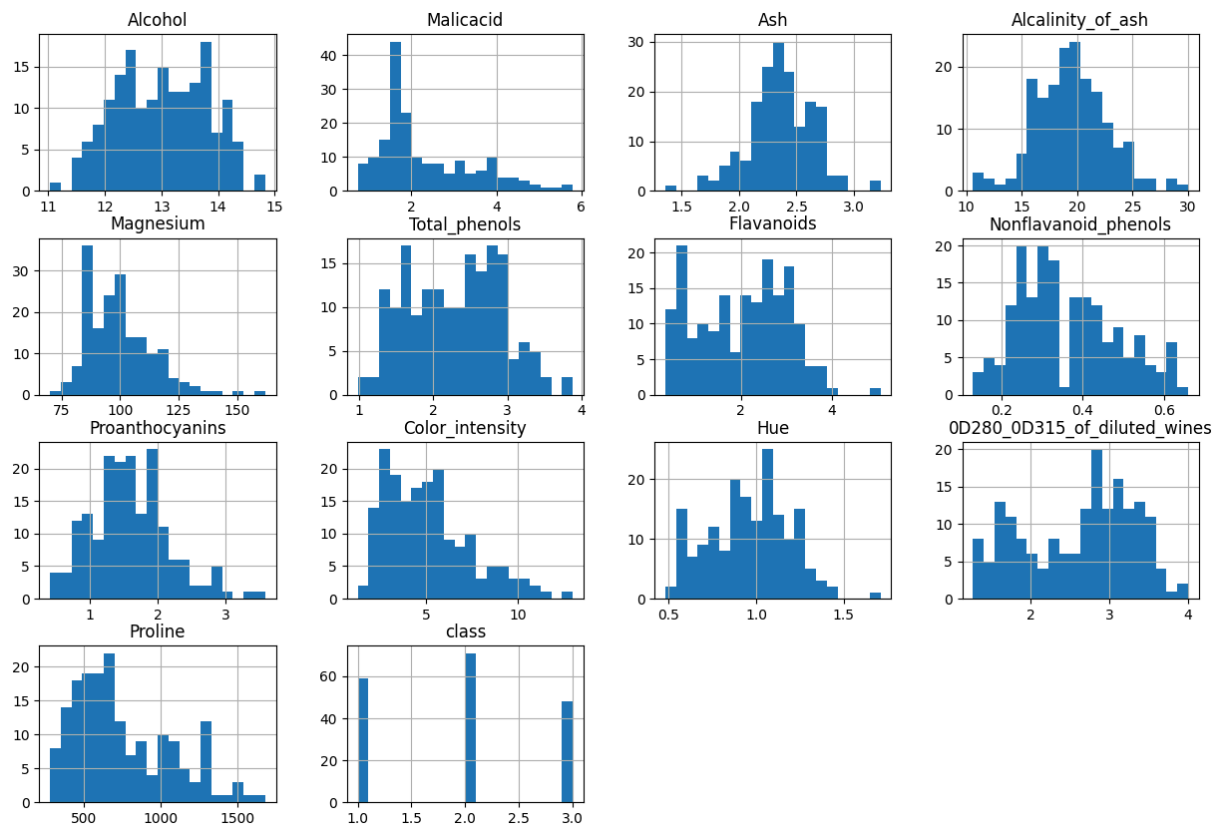
name	Color_intensity	Hue	0D280_0D315_of_diluted_wines	Proline	\
count	178.000000	178.000000	178.000000	178.000000	
mean	5.058090	0.957449	2.611685	746.893258	
std	2.318286	0.228572	0.709990	314.907474	
min	1.280000	0.480000	1.270000	278.000000	
25%	3.220000	0.782500	1.937500	500.500000	
50%	4.690000	0.965000	2.780000	673.500000	
75%	6.200000	1.120000	3.170000	985.000000	
max	13.000000	1.710000	4.000000	1680.000000	

name	class
count	178.000000
mean	1.938202
std	0.775035
min	1.000000
25%	1.000000

```
50%      2.000000
75%      3.000000
max       3.000000
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 178 entries, 0 to 177
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Alcohol                              178 non-null    float64
1   Malicacid                            178 non-null    float64
2   Ash                                  178 non-null    float64
3   Alcalinity_of_ash                    178 non-null    float64
4   Magnesium                            178 non-null    int64
5   Total_phenols                        178 non-null    float64
6   Flavanoids                           178 non-null    float64
7   Nonflavanoid_phenols                 178 non-null    float64
8   Proanthocyanins                      178 non-null    float64
9   Color_intensity                      178 non-null    float64
10  Hue                                  178 non-null    float64
11  0D280_0D315_of_diluted_wines        178 non-null    float64
12  Proline                              178 non-null    int64
13  class                                178 non-null    int64
dtypes: float64(11), int64(3)
memory usage: 19.6 KB
None
```

Data Visualization

```
In [ ]: wine.hist(bins=20, figsize=(15, 10))
plt.show()
```

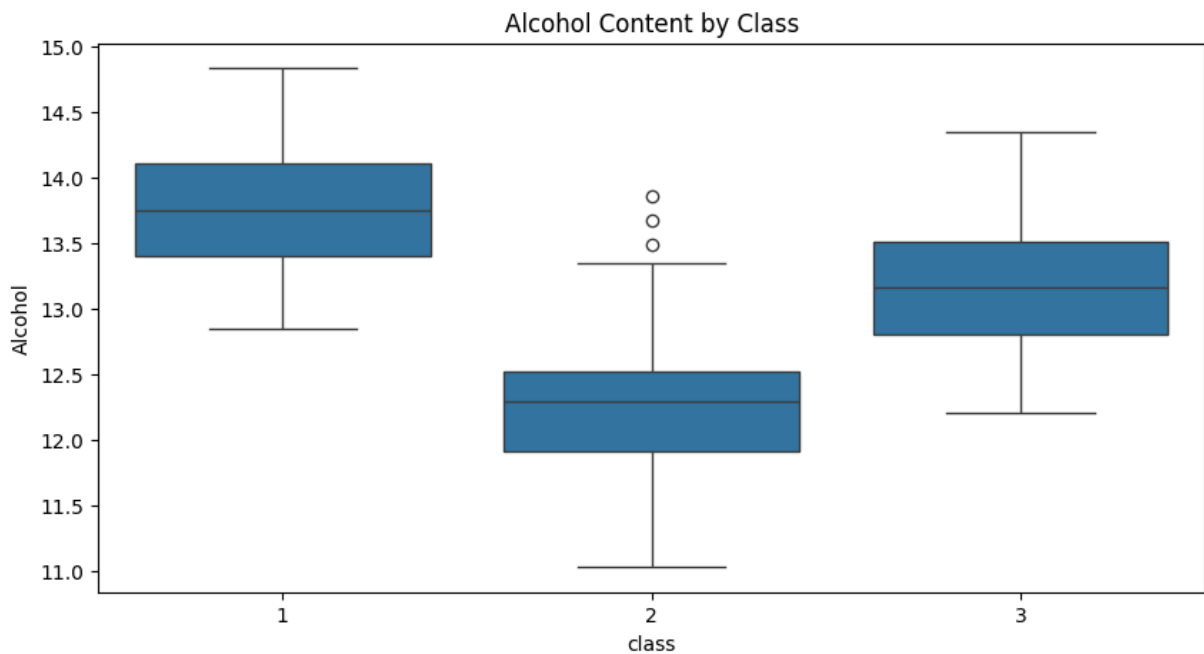


visualize relationships between features

```
In [ ]: sns.pairplot(wine, hue='class')
plt.show()
```



```
In [ ]: plt.figure(figsize=(10, 5))
sns.boxplot(x='class', y='Alcohol', data=wine)
plt.title('Alcohol Content by Class')
plt.show()
```

Data Pre-processing

```
In [ ]: wine.isnull().sum()
```

```
Out[ ]: name
Alcohol                0
Malicacid              0
Ash                   0
Alcalinity_of_ash     0
Magnesium             0
Total_phenols         0
Flavanoids            0
Nonflavanoid_phenols  0
Proanthocyanins       0
Color_intensity       0
Hue                   0
OD280_OD315_of_diluted_wines  0
Proline               0
class                 0
dtype: int64
```


```
In [ ]: scaler = StandardScaler()
wine_scaled = scaler.fit_transform(wine.drop('class', axis=1))
```

```
In [ ]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
wine['class'] = le.fit_transform(wine['class'])
```

```
In [ ]: plt.figure(figsize=(12, 8))
sns.heatmap(wine.corr(), annot=True, cmap='rocket', fmt='.2f')
```

```
plt.title('Correlation Heatmap')
plt.show()
```



Conclusion

Linear Regression is a fundamental technique in data analysis that is used to model the relationship between a dependent variable and one or more independent variables. The aim of linear regression is to find the best line that fits the data, which can be used to make predictions or forecasts. Moreover, Linear regression can be used for a variety of purposes, including predictive modeling, forecasting, exploratory data analysis, and model selection. It is a versatile technique that can be used for a variety of applications, including sales forecasting, stock price predictions, and even weather forecasting. Logistic Regression on the other hand can be use to find answers to questions that have two or more finite outcomes. We can also use logistic regression to pre-process data such as sorting data with a large range of values like bank transactions into a smaller and finite range of values.