

About the data In this notebook, we will using daily weather data that was taken from the National Centers for Environmental Information (NCEI) API. The data collection notebook contains the process that was followed to collect the data. Note: The NCEI is part of the National Oceanic and Atmospheric Administration (NOAA) and, as you can see from the URL for the API, this resource was created when the NCEI was called the NCDC. Should the URL for this resource change in the future, you can search for the NCEI weather API to find the updated one.

Background on the data

Data meanings:

- PRCP : precipitation in millimeters
- SNOW : snowfall in millimeters
- SNWD : snow depth in millimeters
- TMAX : maximum daily temperature in Celsius
- TMIN : minimum daily temperature in Celsius
- TOBS : temperature at time of observation in Celsius
- WESF : water equivalent of snow in millimeters

Setup

```
In [ ]: import pandas as pd
weather = pd.read_csv('/content/nyc_weather_2018.csv')
weather.head()
```

```
Out[ ]:
```

	date	datatype	station	attributes	value
0	2018-01-01T00:00:00	PRCP	GHCND:US1CTFR0039	„N,0800	0.0
1	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0015	„N,1050	0.0
2	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0015	„N,1050	0.0
3	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0017	„N,0920	0.0
4	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0017	„N,0920	0.0

Querying DataFrames

The query() method is an easier way of filtering based on some criteria. For example, we can use it to find all entries where snow was recorded:

```
In [ ]: snow_data = weather.query('datatype == "SNOW" and value > 0')
snow_data.head()

#it gives a query of only the datatype "snow" that has values > 0
```

```
Out[ ]:
```

	date	datatype	station	attributes	value
127	2018-01-01T00:00:00	SNOW	GHCND:US1NYWC0019	„N,1700	25.0
816	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0015	„N,1600	229.0
819	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0017	„N,0830	10.0
823	2018-01-04T00:00:00	SNOW	GHCND:US1NJBG0018	„N,0910	46.0
830	2018-01-04T00:00:00	SNOW	GHCND:US1NYES0018	„N,0700	10.0

This is equivalent to querying the data/weather.db SQLite database for SELECT * FROM weather WHERE datatype == "SNOW" AND value > 0 :

```
In [ ]: import sqlite3
with sqlite3.connect('/content/weather.db') as connection:
    snow_data_from_db = pd.read_sql(
        'SELECT * FROM weather WHERE datatype == "SNOW" AND value > 0',
        connection
    )

snow_data.reset_index().drop(columns='index').equals(snow_data_from_db)

#it's like making a T-SQL in SSMS but in python
```

```
Out[ ]: True
```

Note this is also equivalent to creating Boolean masks:

```
In [ ]: weather[(weather.datatype == 'SNOW') & (weather.value > 0)].equals(snow_data)
```

```
Out[ ]: True
```

Merging DataFrames

We have data for many different stations each day; however, we don't know what the stations are just their IDs. We can join the data in the data/weather_stations.csv file which contains information from the stations endpoint of the NCEI API. Consult the weather_data_collection.ipynb notebook to see how this was collected.

It looks like this:

```
In [ ]: station_info = pd.read_csv('/content/weather_stations.csv')
station_info.head()
```

```
Out[ ]:
```

	id	name	latitude	longitude	elevation
0	GHCND:US1CTFR0022	STAMFORD 2.6 SSW, CT US	41.064100	-73.577000	36.6
1	GHCND:US1CTFR0039	STAMFORD 4.2 S, CT US	41.037788	-73.568176	6.4
2	GHCND:US1NJBG0001	BERGENFIELD 0.3 SW, NJ US	40.921298	-74.001983	20.1
3	GHCND:US1NJBG0002	SADDLE BROOK TWP 0.6 E, NJ US	40.902694	-74.083358	16.8
4	GHCND:US1NJBG0003	TENAFLY 1.3 W, NJ US	40.914670	-73.977500	21.6

As a reminder, the weather data looks like this:

```
In [ ]: weather.head()
```

```
Out[ ]:
```

	date	datatype	station	attributes	value
0	2018-01-01T00:00:00	PRCP	GHCND:US1CTFR0039	„N,0800	0.0
1	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0015	„N,1050	0.0
2	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0015	„N,1050	0.0
3	2018-01-01T00:00:00	PRCP	GHCND:US1NJBG0017	„N,0920	0.0
4	2018-01-01T00:00:00	SNOW	GHCND:US1NJBG0017	„N,0920	0.0

We can join our data by matching up the station_info.id column with the weather.station column. Before doing that though, let's see how many unique values we have:

```
In [ ]: station_info.id.describe()
```

```
Out[ ]: count          320
unique          320
top    GHCND:US1CTFR0022
freq              1
Name: id, dtype: object
```

While station_info has one row per station, the weather dataframe has many entries per station. Notice it also has fewer uniques:

```
In [ ]: weather.station.describe()
```

```
Out[ ]: count          90310
        unique         114
        top      GHCND:USW00014734
        freq          6669
        Name: station, dtype: object
```

When working with joins, it is important to keep an eye on the row count. Some join types will lead to data loss:

```
In [ ]: station_info.shape[0], weather.shape[0]
```

```
Out[ ]: (320, 90310)
```

Since we will be doing this often, it makes more sense to write a function:

```
In [ ]: def get_row_count(*dfs):
        return [df.shape[0] for df in dfs]
        get_row_count(station_info, weather)
```

```
Out[ ]: [320, 90310]
```

The `map()` function is more efficient than list comprehensions. We can couple this with `getattr()` to grab any attribute for multiple dataframes:

```
In [ ]: def get_info(attr, *dfs):
        return list(map(lambda x: getattr(x, attr), dfs))
        get_info('shape', station_info, weather)
```

```
Out[ ]: [(320, 5), (90310, 5)]
```

By default `merge()` performs an inner join. We simply specify the columns to use for the join. The left dataframe is the one we call `merge()` on, and the right one is passed in as an argument:

```
In [ ]: inner_join = weather.merge(station_info, left_on='station', right_on='id')
        inner_join.sample(5, random_state=0)
```

Out[]:

	date	datatype	station	attributes	value	i
51903	2018-04-21T00:00:00	SNOW	GHCND:USW00014734	„W,	0.0	GHCND:USW0001473
11216	2018-11-13T00:00:00	DAPR	GHCND:US1NJMN0104	„N,2359	5.0	GHCND:US1NJMN010
13292	2018-04-02T00:00:00	PRCP	GHCND:US1NJMS0059	„N,2200	15.2	GHCND:US1NJMS005
51540	2018-04-02T00:00:00	AWBT	GHCND:USW00014734	„W,	11.0	GHCND:USW0001473
66282	2018-07-09T00:00:00	RHAV	GHCND:USW00094728	„W,	51.0	GHCND:USW0009472

We can remove the duplication of information in the station and id columns by renaming one of them before the merge and then simply using on :

```
In [ ]: weather.merge(station_info.rename(dict(id='station'), axis=1), on='station').sample
```

Out[]:

	date	datatype	station	attributes	value	name	li
51903	2018-04-21T00:00:00	SNOW	GHCND:USW00014734	„W,	0.0	NEWARK LIBERTY INTERNATIONAL AIRPORT, NJ US	40.
11216	2018-11-13T00:00:00	DAPR	GHCND:US1NJMN0104	„N,2359	5.0	LITTLE SILVER 0.3 NNW, NJ US	40.
13292	2018-04-02T00:00:00	PRCP	GHCND:US1NJMS0059	„N,2200	15.2	MADISON 0.8 WSW, NJ US	40.
51540	2018-04-02T00:00:00	AWBT	GHCND:USW00014734	„W,	11.0	NEWARK LIBERTY INTERNATIONAL AIRPORT, NJ US	40.
66282	2018-07-09T00:00:00	RHAV	GHCND:USW00094728	„W,	51.0	NY CITY CENTRAL PARK, NY US	40.

We are losing stations that don't have weather observations associated with them, if we don't want to lose these rows, we perform a right or left join instead of the inner join:

```
In [ ]: left_join = station_info.merge(weather, left_on='id', right_on='station', how='left')
right_join = weather.merge(station_info, left_on='station', right_on='id', how='right')
right_join.tail()
```

```
Out[ ]:
```

	date	datatype	station	attributes	value	id
90511	2018-12-31T00:00:00	WDF5	GHCND:USW00094789	„W,	130.0	GHCND:USW00094789
90512	2018-12-31T00:00:00	WSF2	GHCND:USW00094789	„W,	9.8	GHCND:USW00094789
90513	2018-12-31T00:00:00	WSF5	GHCND:USW00094789	„W,	12.5	GHCND:USW00094789
90514	2018-12-31T00:00:00	WT01	GHCND:USW00094789	„W,	1.0	GHCND:USW00094789
90515	2018-12-31T00:00:00	WT02	GHCND:USW00094789	„W,	1.0	GHCND:USW00094789

The left and right join as we performed above are equivalent because the side that we kept the rows without matches was the same in both cases:

```
In [ ]: left_join.sort_index(axis=1).sort_values(['date', 'station']).reset_index().drop(columns='id')
right_join.sort_index(axis=1).sort_values(['date', 'station']).reset_index().drop(columns='id')
```

```
Out[ ]: True
```

Note we have additional rows in the left and right joins because we kept all the stations that didn't have weather observations:

```
In [ ]: get_info('shape', inner_join, left_join, right_join)
```

```
Out[ ]: [(90310, 10), (90516, 10), (90516, 10)]
```

If we query the station information for stations that have NY in their name, believing that to be all the stations that record weather data for NYC and perform an outer join, we can see where the mismatches occur:

```
In [ ]: outer_join = weather.merge(
    station_info[station_info.name.str.contains('NY')],
    left_on='station', right_on='id', how='outer', indicator=True
)
```

```
outer_join.sample(4, random_state=0).append(outer_join[outer_join.station.isna()]).h
```

<ipython-input-23-17f7c35892b9>:6: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
outer_join.sample(4, random_state=0).append(outer_join[outer_join.station.isna()]).head(2))
```

Out []:

	date	datatype	station	attributes	value	ic
60709	2018-04-15T00:00:00	WSF5	GHCND:USW00054787	„W,	15.2	GHCND:USW00054787
49486	2018-12-09T00:00:00	WDF5	GHCND:USW00014732	„W,	310.0	GHCND:USW00014732
3884	2018-09-30T00:00:00	PRCP	GHCND:US1NJES0018	„N,0700	0.0	NaN
22869	2018-12-30T00:00:00	PRCP	GHCND:US1NYNS0030	„N,0800	0.0	GHCND:US1NYNS0030
90310	NaN	NaN	NaN	NaN	NaN	GHCND:US1NJHD0018
90311	NaN	NaN	NaN	NaN	NaN	GHCND:US1NJMS0036

These joins are equivalent to their SQL counterparts. Below is the inner join. Note that to use equals() you will have to do some manipulation of the dataframes to line them up:

```
In [ ]: import sqlite3
with sqlite3.connect('/content/weather.db') as connection:
    inner_join_from_db = pd.read_sql(
        'SELECT * FROM weather JOIN stations ON weather.station == stations.id',
        connection
    )

inner_join_from_db.shape == inner_join.shape
```

Out []: True

Revisit the dirty data from the previous module.

```
In [ ]: dirty_data = pd.read_csv(
    '/content/dirty_data.csv', index_col='date'
).drop_duplicates().drop(columns='SNWD')
dirty_data.head()
```

Out[]:

	station	PRCP	SNOW	TMAX	TMIN	TOBS	WESF	inclement_w
date								
2018-01-01T00:00:00	?	0.0	0.0	5505.0	-40.0	NaN	NaN	
2018-01-02T00:00:00	GHCND:USC00280907	0.0	0.0	-8.3	-16.1	-12.2	NaN	
2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-4.4	-13.9	-13.3	NaN	
2018-01-04T00:00:00	?	20.6	229.0	5505.0	-40.0	NaN	19.3	
2018-01-05T00:00:00	?	0.3	NaN	5505.0	-40.0	NaN	NaN	

We need to create two dataframes for the join. We will drop some unnecessary columns as well for easier viewing:

```
In [ ]: valid_station = dirty_data.query('station != "?"').copy().drop(columns=['WESF', 'st
station_with_wesf = dirty_data.query('station == "?"').copy().drop(columns=['statio
```

Our column for the join is the index in both dataframes, so we must specify left_index and right_index :

```
In [ ]: valid_station.merge(
    station_with_wesf, left_index=True, right_index=True
).query('WESF > 0').head()
```

Out[]:

	PRCP_x	SNOW_x	TMAX	TMIN	TOBS	inclement_weather_x	PRCP_y	SNOW
date								
2018-01-30T00:00:00	0.0	0.0	6.7	-1.7	-0.6	False	1.5	1.
2018-03-08T00:00:00	48.8	NaN	1.1	-0.6	1.1	False	28.4	N.
2018-03-13T00:00:00	4.1	51.0	5.6	-3.9	0.0	True	3.0	1.
2018-03-21T00:00:00	0.0	0.0	2.8	-2.8	0.6	False	6.6	11.
2018-04-02T00:00:00	9.1	127.0	12.8	-1.1	-1.1	True	14.0	15.

The columns that existed in both dataframes, but didn't form part of the join got suffixes added to their names: `_x` for columns from the left dataframe and `_y` for columns from the right dataframe. We can customize this with the `suffixes` argument:

```
In [ ]: valid_station.merge(
        station_with_wesf, left_index=True, right_index=True, suffixes=('_', '_?'))
        ).query('WESF > 0').head()
```

```
Out[ ]:      PRCP  SNOW  TMAX  TMIN  TOBS  inclement_weather  PRCP_?  SNOW_?  W
date
```

2018-01-30T00:00:00	0.0	0.0	6.7	-1.7	-0.6	False	1.5	13.0
2018-03-08T00:00:00	48.8	NaN	1.1	-0.6	1.1	False	28.4	NaN
2018-03-13T00:00:00	4.1	51.0	5.6	-3.9	0.0	True	3.0	13.0
2018-03-21T00:00:00	0.0	0.0	2.8	-2.8	0.6	False	6.6	114.0
2018-04-02T00:00:00	9.1	127.0	12.8	-1.1	-1.1	True	14.0	152.0

Since we are joining on the index, an easier way is to use the `join()` method instead of `merge()`. Note that the suffix parameter is now `lsuffix` for the left dataframe's suffix and `rsuffix` for the right one's:

```
In [ ]: valid_station.join(station_with_wesf, rsuffix='_?').query('WESF > 0').head()
```

```
Out[ ]:      PRCP  SNOW  TMAX  TMIN  TOBS  inclement_weather  PRCP_?  SNOW_?  W
date
```

2018-01-30T00:00:00	0.0	0.0	6.7	-1.7	-0.6	False	1.5	13.0
2018-03-08T00:00:00	48.8	NaN	1.1	-0.6	1.1	False	28.4	NaN
2018-03-13T00:00:00	4.1	51.0	5.6	-3.9	0.0	True	3.0	13.0
2018-03-21T00:00:00	0.0	0.0	2.8	-2.8	0.6	False	6.6	114.0
2018-04-02T00:00:00	9.1	127.0	12.8	-1.1	-1.1	True	14.0	152.0

Joins can be very resource-intensive, so it's a good idea to figure out what type of join you need using set operations before trying the join itself. The pandas set operations are performed on the index, so whichever columns we will be joining on will need to be the index. Let's go back to the weather and station_info dataframes and set the station ID columns as the index:

```
In [ ]: weather.set_index('station', inplace=True)
        station_info.set_index('id', inplace=True)
```

The intersection will tell us the stations that are present in both dataframes. The result will be the index when performing an inner join:

```
In [ ]: weather.index.intersection(station_info.index)
```

```
Out[ ]: Index(['GHCND:US1CTFR0039', 'GHCND:US1NJBG0015', 'GHCND:US1NJBG0017',
              'GHCND:US1NJBG0018', 'GHCND:US1NJBG0023', 'GHCND:US1NJBG0030',
              'GHCND:US1NJBG0039', 'GHCND:US1NJBG0044', 'GHCND:US1NJBG0018',
              'GHCND:US1NJBG0024',
              ...
              'GHCND:USC00284987', 'GHCND:US1NJBG0031', 'GHCND:US1NJBG0029',
              'GHCND:US1NJBG0086', 'GHCND:US1NJBG0097', 'GHCND:US1NJBG0081',
              'GHCND:US1NJBG0088', 'GHCND:US1NJBG0033', 'GHCND:US1NJBG0040',
              'GHCND:US1NJBG0029'],
              dtype='object', length=114)
```

The set difference will tell us what we lose from each side. When performing an inner join, we lose nothing from the weather dataframe:

```
In [ ]: weather.index.difference(station_info.index)
```

```
Out[ ]: Index([], dtype='object')
```

We lose 153 stations from the station_info dataframe, however:

```
In [ ]: station_info.index.difference(weather.index)
```

```
Out[ ]: Index(['GHCND:US1CTFR0022', 'GHCND:US1NJBG0001', 'GHCND:US1NJBG0002',
              'GHCND:US1NJBG0005', 'GHCND:US1NJBG0006', 'GHCND:US1NJBG0008',
              'GHCND:US1NJBG0011', 'GHCND:US1NJBG0012', 'GHCND:US1NJBG0013',
              'GHCND:US1NJBG0020',
              ...
              'GHCND:USC00308749', 'GHCND:USC00308946', 'GHCND:USC00309117',
              'GHCND:USC00309270', 'GHCND:USC00309400', 'GHCND:USC00309466',
              'GHCND:USC00309576', 'GHCND:USC00309580', 'GHCND:USW00014708',
              'GHCND:USW00014786'],
              dtype='object', length=206)
```

The symmetric difference will tell us what gets lost from both sides. It is the combination of the set difference in both directions:

```
In [ ]: ny_in_name = station_info[station_info.name.str.contains('NY')]

ny_in_name.index.difference(weather.index).shape[0]\
+ weather.index.difference(ny_in_name.index).shape[0]\
== weather.index.symmetric_difference(ny_in_name.index).shape[0]
```

```
Out[ ]: True
```

The union will show us everything that will be present after a full outer join. Note that since these are sets (which don't allow duplicates by definition), we must pass unique entries for union:

```
In [ ]: weather.index.unique().union(station_info.index)
```

```
Out[ ]: Index(['GHCND:US1CTFR0022', 'GHCND:US1CTFR0039', 'GHCND:US1NJBG0001',
              'GHCND:US1NJBG0002', 'GHCND:US1NJBG0003', 'GHCND:US1NJBG0005',
              'GHCND:US1NJBG0006', 'GHCND:US1NJBG0008', 'GHCND:US1NJBG0010',
              'GHCND:US1NJBG0011',
              ...
              'GHCND:USW00014708', 'GHCND:USW00014732', 'GHCND:USW00014734',
              'GHCND:USW00014786', 'GHCND:USW00054743', 'GHCND:USW00054787',
              'GHCND:USW00094728', 'GHCND:USW00094741', 'GHCND:USW00094745',
              'GHCND:USW00094789'],
              dtype='object', length=320)
```

Note that the symmetric difference is actually the union of the set differences:

```
In [ ]: ny_in_name = station_info[station_info.name.str.contains('NY')]

ny_in_name.index.difference(weather.index).union(weather.index.difference(ny_in_name.index)
          weather.index.symmetric_difference(ny_in_name.index)
          )
```

```
Out[ ]: True
```