

Time Series

Setup

```
In [1]: import numpy as np
import pandas as pd

fb = pd.read_csv('/content/fb_2018.csv', index_col='date', parse_dates=True).assign(
    trading_volume=lambda x: pd.cut(x.volume, bins=3, labels=['low', 'med', 'high'])
)
fb.head()
```

```
Out[1]:
```

	open	high	low	close	volume	trading_volume
date						
2018-01-02	177.68	181.58	177.5500	181.42	18151903	low
2018-01-03	181.88	184.78	181.3300	184.67	16886563	low
2018-01-04	184.90	186.21	184.0996	184.33	13880896	low
2018-01-05	185.59	186.90	184.9300	186.85	13574535	low
2018-01-08	187.20	188.90	186.3300	188.28	17994726	low

Time-based selection and filtering

Remember, when we have a DatetimeIndex, we can use datetime slicing. We can provide a range of dates. We only get three days back because the stock market is closed on the weekends:

```
In [2]: fb['2018-10-11':'2018-10-15']

#through index slicing we were able to get the range of dates
```

```
Out[2]:
```

	open	high	low	close	volume	trading_volume
date						
2018-10-11	150.13	154.81	149.1600	153.35	35338901	low
2018-10-12	156.73	156.89	151.2998	153.74	25293492	low
2018-10-15	153.32	155.57	152.5500	153.52	15433521	low

We can select ranges of months and quarters:

```
In [3]: fb['2018-q1'].equals(fb['2018-01':'2018-03'])
```

```
<ipython-input-3-f01e3c270a70>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the rows, like `frame[string]`, is deprecated and will be removed in a future version. Use `frame.loc[string]` instead.
fb['2018-q1'].equals(fb['2018-01':'2018-03'])
```

```
Out[3]: True
```

The first() method will give us a specified length of time from the beginning of the time series. Here, we ask for a week. January 1, 2018 was a holiday—meaning the market was closed. It was also a Monday, so the week here is only four days:

```
In [4]: fb.first('1W')
#this gives us an output where it shows us the data of the first week of the stock
```

```
Out[4]:
```

	open	high	low	close	volume	trading_volume
date						
2018-01-02	177.68	181.58	177.5500	181.42	18151903	low
2018-01-03	181.88	184.78	181.3300	184.67	16886563	low
2018-01-04	184.90	186.21	184.0996	184.33	13880896	low
2018-01-05	185.59	186.90	184.9300	186.85	13574535	low

The last() method will take from the end:

```
In [5]: fb.last('1W')
#the output gives us the data of the last week of the stock market
```

```
Out[5]:
```

	open	high	low	close	volume	trading_volume
date						
2018-12-31	134.45	134.64	129.95	131.09	24625308	low

For the next few examples, we need datetimes, so we will read in the stock data per minute file:

```
In [7]: stock_data_per_minute = pd.read_csv(
    '/content/fb_week_of_may_20_per_minute.csv', index_col='date', parse_dates=True,
    date_parser=lambda x: pd.to_datetime(x, format='%Y-%m-%d %H-%M')
)

stock_data_per_minute.head()
```

Out[7]:

	open	high	low	close	volume
date					
2019-05-20 09:30:00	181.6200	181.6200	181.6200	181.6200	159049.0
2019-05-20 09:31:00	182.6100	182.6100	182.6100	182.6100	468017.0
2019-05-20 09:32:00	182.7458	182.7458	182.7458	182.7458	97258.0
2019-05-20 09:33:00	182.9500	182.9500	182.9500	182.9500	43961.0
2019-05-20 09:34:00	183.0600	183.0600	183.0600	183.0600	79562.0

We can use the Grouper to roll up our data to the daily level along with first and last :

```
In [8]: stock_data_per_minute.groupby(pd.Grouper(freq='1D')).agg({
    'open': 'first',
    'high': 'max',
    'low': 'min',
    'close': 'last',
    'volume': 'sum'
})
```

Out[8]:

	open	high	low	close	volume
date					
2019-05-20	181.62	184.1800	181.6200	182.72	10044838.0
2019-05-21	184.53	185.5800	183.9700	184.82	7198405.0
2019-05-22	184.81	186.5603	184.0120	185.32	8412433.0
2019-05-23	182.50	183.7300	179.7559	180.87	12479171.0
2019-05-24	182.33	183.5227	181.0400	181.06	7686030.0

The `at_time()` method allows us to pull out all datetimes that match a certain time. Here, we can grab all the rows from the time the stock market opens (9:30 AM):

```
In [10]: stock_data_per_minute.at_time('9:30')
```

```
Out[10]:
```

	open	high	low	close	volume
date					
2019-05-20 09:30:00	181.62	181.62	181.62	181.62	159049.0
2019-05-21 09:30:00	184.53	184.53	184.53	184.53	58171.0
2019-05-22 09:30:00	184.81	184.81	184.81	184.81	41585.0
2019-05-23 09:30:00	182.50	182.50	182.50	182.50	121930.0
2019-05-24 09:30:00	182.33	182.33	182.33	182.33	52681.0

We can use `between_time()` to grab data for the last two minutes of trading daily:

```
In [11]: stock_data_per_minute.between_time('15:59', '16:00')
```

```
Out[11]:
```

	open	high	low	close	volume
date					
2019-05-20 15:59:00	182.915	182.915	182.915	182.915	134569.0
2019-05-20 16:00:00	182.720	182.720	182.720	182.720	1113672.0
2019-05-21 15:59:00	184.840	184.840	184.840	184.840	61606.0
2019-05-21 16:00:00	184.820	184.820	184.820	184.820	801080.0
2019-05-22 15:59:00	185.290	185.290	185.290	185.290	96099.0
2019-05-22 16:00:00	185.320	185.320	185.320	185.320	1220993.0
2019-05-23 15:59:00	180.720	180.720	180.720	180.720	109648.0
2019-05-23 16:00:00	180.870	180.870	180.870	180.870	1329217.0
2019-05-24 15:59:00	181.070	181.070	181.070	181.070	52994.0
2019-05-24 16:00:00	181.060	181.060	181.060	181.060	764906.0

On average, are more shares traded within the first 30 minutes of trading or in the last 30 minutes? We can combine `between_time()` with `Groupers` and `filter()` from the `aggregation.ipynb` notebook to answer this question. For the week in question, more are traded on average around opening time than closing time:

```
In [12]: shares_traded_in_first_30_min = stock_data_per_minute\
        .between_time('9:30', '10:00')\
        .groupby(pd.Grouper(freq='1D'))\
        .filter(lambda x: (x.volume > 0).all())\
        .volume.mean()

shares_traded_in_last_30_min = stock_data_per_minute\
        .between_time('15:30', '16:00')
```

```
.groupby(pd.Grouper(freq='1D'))\
.filter(lambda x: (x.volume > 0).all())\
.volume.mean()
```

```
shares_traded_in_first_30_min - shares_traded_in_last_30_min
```

Out[12]: 18592.967741935485

In cases where time doesn't matter, we can normalize the times to midnight:

```
In [13]: pd.DataFrame(
    dict(before=stock_data_per_minute.index, after=stock_data_per_minute.index.normalize()).head()
```

Out[13]:

	before	after
0	2019-05-20 09:30:00	2019-05-20
1	2019-05-20 09:31:00	2019-05-20
2	2019-05-20 09:32:00	2019-05-20
3	2019-05-20 09:33:00	2019-05-20
4	2019-05-20 09:34:00	2019-05-20

Note that we can also use `normalize()` on a Series object after accessing the `dt` attribute:

```
In [14]: stock_data_per_minute.index.to_series().dt.normalize().head()
```

Out[14]:

```
date
2019-05-20 09:30:00    2019-05-20
2019-05-20 09:31:00    2019-05-20
2019-05-20 09:32:00    2019-05-20
2019-05-20 09:33:00    2019-05-20
2019-05-20 09:34:00    2019-05-20
Name: date, dtype: datetime64[ns]
```

Shifting for lagged data

We can use `shift()` to create some lagged data. By default, the shift will be one period. For example, we can use `shift()` to create a new column that indicates the previous day's closing price. From this new column, we can calculate the price change due to after hours trading (after the close one day right up to the open the following day):

```
In [15]: fb.assign(
    prior_close=lambda x: x.close.shift(),
    after_hours_change_in_price=lambda x: x.open - x.prior_close,
    abs_change=lambda x: x.after_hours_change_in_price.abs()
).nlargest(5, 'abs_change')
```

Out[15]:

	open	high	low	close	volume	trading_volume	prior_close	after_hours_
date								
2018-07-26	174.89	180.13	173.75	176.26	169803668	high	217.50	
2018-04-26	173.22	176.27	170.80	174.16	77556934	med	159.69	
2018-01-12	178.06	181.48	177.40	179.37	77551299	med	187.77	
2018-10-31	155.00	156.40	148.96	151.79	60101251	low	146.22	
2018-03-19	177.01	177.17	170.06	172.56	88140060	med	185.09	

The `tshift()` method will shift the `DatetimeIndex` rather than the data. However, if the goal is to add/subtract time we can use `pd.Timedelta`:

In [16]: `pd.date_range('2018-01-01', freq='D', periods=5) + pd.Timedelta('9 hours 30 minutes')`

Out[16]: `DatetimeIndex(['2018-01-01 09:30:00', '2018-01-02 09:30:00',
'2018-01-03 09:30:00', '2018-01-04 09:30:00',
'2018-01-05 09:30:00'],
dtype='datetime64[ns]', freq='D')`

When working with stock data, we only have data for the dates the market was open. We can use `first_valid_index()` to give us the index of the first non-null entry in our data. For September 2018, this is September 4th:

In [17]: `fb['2018-09'].first_valid_index()`

<ipython-input-17-d8ca41528993>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the rows, like `frame[string]`, is deprecated and will be removed in a future version. Use `frame.loc[string]` instead.
`fb['2018-09'].first_valid_index()`

Out[17]: `Timestamp('2018-09-04 00:00:00')`

Conversely, we can use `last_valid_index()` to get the last entry of non-null data. For September 2018, this is September 28th:

In [18]: `fb['2018-09'].last_valid_index()`

<ipython-input-18-ef6e024573c9>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the rows, like `frame[string]`, is deprecated and will be removed in a future version. Use `frame.loc[string]` instead.
`fb['2018-09'].last_valid_index()`

Out[18]: `Timestamp('2018-09-28 00:00:00')`

We can use `asof()` to find the last non-null data before the point we are looking for, if it isn't in the index. From the previous result, we know that the market was not open on September 30th. It also isn't in the index:

```
In [20]: fb.index.contains('2018-09-30')
#by looking through the documentation of index.contains, it should be in __contains
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-20-57e448819eff> in <cell line: 1>()
----> 1 fb.index.contains('2018-09-30')

AttributeError: 'DatetimeIndex' object has no attribute 'contain'
```

```
In [45]: fb.index.__contains__('2018-09-30')
```

```
Out[45]: False
```

If we ask for it, we will get the data from the index we got from `fb['2018-09'].last_valid_index()`, which was September 28th:

```
In [21]: fb.asof('2018-09-30')
```

```
Out[21]: open                168.33
         high                168.79
         low                 162.56
         close               164.46
         volume             34265638
         trading_volume      low
         Name: 2018-09-30 00:00:00, dtype: object
```

Differenced data

Using the `diff()` method is a quick way to calculate the difference between the data and a lagged version of it. By default, it will yield the result of `data - data.shift()`:

```
In [22]: (
         fb.drop(columns='trading_volume')
         - fb.drop(columns='trading_volume').shift()
         ).equals(
         fb.drop(columns='trading_volume').diff()
         )
```

```
Out[22]: True
```

We can use this to see how Facebook stock changed day-over-day:

```
In [23]: fb.drop(columns='trading_volume').diff().head()
```

Out[23]:

	open	high	low	close	volume
date					
2018-01-02	NaN	NaN	NaN	NaN	NaN
2018-01-03	4.20	3.20	3.7800	3.25	-1265340.0
2018-01-04	3.02	1.43	2.7696	-0.34	-3005667.0
2018-01-05	0.69	0.69	0.8304	2.52	-306361.0
2018-01-08	1.61	2.00	1.4000	1.43	4420191.0

We can specify the number of periods, can be any positive or negative integer:

In [24]: `fb.drop(columns='trading_volume').diff(-3).head()`

Out[24]:

	open	high	low	close	volume
date					
2018-01-02	-7.91	-5.32	-7.3800	-5.43	4577368.0
2018-01-03	-5.32	-4.12	-5.0000	-3.61	-1108163.0
2018-01-04	-3.80	-2.59	-3.0004	-3.54	1487839.0
2018-01-05	-1.35	-0.99	-0.7000	-0.99	3044641.0
2018-01-08	-1.20	0.50	-1.0500	0.51	8406139.0

Resampling

Sometimes the data is at a granularity that isn't conducive to our analysis. Consider the case where we have data per minute for the full year of 2018. Let's see what happens if we try to plot this.

Plotting will be covered in the next module, so don't worry too much about the code.

First, we import matplotlib for plotting:


In [25]: `import matplotlib.pyplot as plt`

Then we will look at the plot at the minute level and at the daily aggregated level (summed):

In [26]: `np.random.seed(0)
index = pd.date_range('2018-01-01', freq='T', periods=365*24*60)
raw = pd.DataFrame(
 np.random.uniform(0, 10, size=index.shape[0]), index=index
)`


```
fig, axes = plt.subplots(1, 2, figsize=(15, 5))
raw.plot(legend=False, ax=axes[0], title='raw data')
raw.resample('1D').sum().plot(legend=False, ax=axes[1], title='daily totals')
for ax in axes:
    ax.set_xlabel('date')
    ax.set_ylabel('events')

plt.suptitle('Raw versus Resampled Data')
plt.show()
```



The plot on the left has so much data we can't see anything. However, when we aggregate to the daily totals, we see the data. We can alter the granularity of the data we are working with using resampling. Recall our minute-by-minute stock data:

In [27]: `stock_data_per_minute.head()`

Out[27]:

	open	high	low	close	volume
date					
2019-05-20 09:30:00	181.6200	181.6200	181.6200	181.6200	159049.0
2019-05-20 09:31:00	182.6100	182.6100	182.6100	182.6100	468017.0
2019-05-20 09:32:00	182.7458	182.7458	182.7458	182.7458	97258.0
2019-05-20 09:33:00	182.9500	182.9500	182.9500	182.9500	43961.0
2019-05-20 09:34:00	183.0600	183.0600	183.0600	183.0600	79562.0

We can resample this to get to a daily frequency:

In [28]: `stock_data_per_minute.resample('1D').agg({`
`'open': 'first',`
`'high': 'max',`
`'low': 'min',`
`'close': 'last',`

```
'volume': 'sum'
})
```

Out[28]:

	open	high	low	close	volume
date					
2019-05-20	181.62	184.1800	181.6200	182.72	10044838.0
2019-05-21	184.53	185.5800	183.9700	184.82	7198405.0
2019-05-22	184.81	186.5603	184.0120	185.32	8412433.0
2019-05-23	182.50	183.7300	179.7559	180.87	12479171.0
2019-05-24	182.33	183.5227	181.0400	181.06	7686030.0

We can downsample to quarterly data:

In [29]: `fb.resample('Q').mean()`

<ipython-input-29-f6fd3d834d43>:1: FutureWarning: The default value of numeric_only in DataFrameGroupBy.mean is deprecated. In a future version, numeric_only will default to False. Either specify numeric_only or select only columns which should be valid for the function.

```
fb.resample('Q').mean()
```

Out[29]:

	open	high	low	close	volume
date					
2018-03-31	179.472295	181.794659	177.040428	179.551148	3.292640e+07
2018-06-30	180.373770	182.277689	178.595964	180.704688	2.405532e+07
2018-09-30	180.812130	182.890886	178.955229	181.028492	2.701982e+07
2018-12-31	145.272460	147.620121	142.718943	144.868730	2.697433e+07

We can also use `apply()`. Here, we show the quarterly change from start to end:

In [30]: `fb.drop(columns='trading_volume').resample('Q').apply(
 lambda x: x.last('1D').values - x.first('1D').values
)`

Out[30]:

date	
2018-03-31	[[-22.53, -20.160000000000025, -23.4100000000...
2018-06-30	[[-39.509999999999999, 38.399700000000024, 39.84...
2018-09-30	[[-25.039999999999999, -28.659999999999997, -2...
2018-12-31	[[-28.580000000000013, -31.24000000000001, -31...

Freq: Q-DEC, dtype: object

Consider the following melted stock data by the minute. We don't see the OHLC data directly:

```
In [32]: melted_stock_data = pd.read_csv('/content/melted_stock_data.csv', index_col='date',
      melted_stock_data.head()
```

```
Out[32]:
```

	price
date	
2019-05-20 09:30:00	181.6200
2019-05-20 09:31:00	182.6100
2019-05-20 09:32:00	182.7458
2019-05-20 09:33:00	182.9500
2019-05-20 09:34:00	183.0600

We can use the ohlc() method after resampling to recover the OHLC columns:

```
In [33]: melted_stock_data.resample('1D').ohlc()['price']
```

```
Out[33]:
```

	open	high	low	close
date				
2019-05-20	181.62	184.1800	181.6200	182.72
2019-05-21	184.53	185.5800	183.9700	184.82
2019-05-22	184.81	186.5603	184.0120	185.32
2019-05-23	182.50	183.7300	179.7559	180.87
2019-05-24	182.33	183.5227	181.0400	181.06

Alternatively, we can upsample to increase the granularity. Note this will introduce NaN values:

```
In [34]: fb.resample('6H').asfreq().head()
```

```
Out[34]:
```

	open	high	low	close	volume	trading_volume
date						
2018-01-02 00:00:00	177.68	181.58	177.55	181.42	18151903.0	low
2018-01-02 06:00:00	NaN	NaN	NaN	NaN	NaN	NaN
2018-01-02 12:00:00	NaN	NaN	NaN	NaN	NaN	NaN
2018-01-02 18:00:00	NaN	NaN	NaN	NaN	NaN	NaN
2018-01-03 00:00:00	181.88	184.78	181.33	184.67	16886563.0	low

There are many ways to handle these NaN values. We can forward-fill with `pad()` :

```
In [35]: fb.resample('6H').pad().head()
```

<ipython-input-35-39179f05e435>:1: FutureWarning: pad is deprecated and will be removed in a future version. Use ffill instead.

```
fb.resample('6H').pad().head()
```

```
Out[35]:
```

	open	high	low	close	volume	trading_volume
date						
2018-01-02 00:00:00	177.68	181.58	177.55	181.42	18151903	low
2018-01-02 06:00:00	177.68	181.58	177.55	181.42	18151903	low
2018-01-02 12:00:00	177.68	181.58	177.55	181.42	18151903	low
2018-01-02 18:00:00	177.68	181.58	177.55	181.42	18151903	low
2018-01-03 00:00:00	181.88	184.78	181.33	184.67	16886563	low

We can specify a specific value or a method with `fillna()` :

```
In [36]: fb.resample('6H').fillna('nearest').head()
```

```
Out[36]:
```

	open	high	low	close	volume	trading_volume
date						
2018-01-02 00:00:00	177.68	181.58	177.55	181.42	18151903	low
2018-01-02 06:00:00	177.68	181.58	177.55	181.42	18151903	low
2018-01-02 12:00:00	181.88	184.78	181.33	184.67	16886563	low
2018-01-02 18:00:00	181.88	184.78	181.33	184.67	16886563	low
2018-01-03 00:00:00	181.88	184.78	181.33	184.67	16886563	low

We can use `asfreq()` and `assign()` to specify the action per column:

```
In [37]: fb.resample('6H').asfreq().assign(
    volume=lambda x: x.volume.fillna(0), # put 0 when market is closed
    close=lambda x: x.close.fillna(method='ffill'), # carry forward
    # take the closing price if these aren't available
    open=lambda x: np.where(x.open.isnull(), x.close, x.open),
    high=lambda x: np.where(x.high.isnull(), x.close, x.high),
    low=lambda x: np.where(x.low.isnull(), x.close, x.low)
).head()
```

Out[37]:

	open	high	low	close	volume	trading_volume
date						
2018-01-02 00:00:00	177.68	181.58	177.55	181.42	18151903.0	low
2018-01-02 06:00:00	181.42	181.42	181.42	181.42	0.0	NaN
2018-01-02 12:00:00	181.42	181.42	181.42	181.42	0.0	NaN
2018-01-02 18:00:00	181.42	181.42	181.42	181.42	0.0	NaN
2018-01-03 00:00:00	181.88	184.78	181.33	184.67	16886563.0	low

Merging

We saw merging examples the querying_and_merging notebook. However, they all matched based on keys. With time series, it is possible that they are so granular that we never have the same time for multiple entries. Let's work with some stock data at different granularities

```
In [39]: import sqlite3
with sqlite3.connect('/content/stocks.db') as connection:
    fb_prices = pd.read_sql(
        'SELECT * FROM fb_prices', connection,
        index_col='date', parse_dates=['date']
    )
    aapl_prices = pd.read_sql(
        'SELECT * FROM aapl_prices', connection,
        index_col='date', parse_dates=['date']
    )
```

The Facebook prices are at the minute granularity:

```
In [40]: fb_prices.index.second.unique()
```

```
Out[40]: Int64Index([0], dtype='int64', name='date')
```

However, the Apple prices have information for the second:

```
In [41]: aapl_prices.index.second.unique()
```

```
Out[41]: Int64Index([ 0, 52, 36, 34, 55, 35,  7, 12, 59, 17,  5, 20, 26, 23, 54, 49, 19,
                    53, 11, 22, 13, 21, 10, 46, 42, 38, 33, 18, 16,  9, 56, 39,  2, 50,
                    31, 58, 48, 24, 29,  6, 47, 51, 40,  3, 15, 14, 25,  4, 43,  8, 32,
                    27, 30, 45,  1, 44, 57, 41, 37, 28],
                    dtype='int64', name='date')
```

We can perform an asof merge to try to line these up the best we can. We specify how to handle the mismatch with the direction and tolerance parameters. We will fill in with the direction of nearest and a tolerance of 30 seconds. This will place the Apple data with the

minute that it is closest to, so 9:31:52 will go with 9:32 and 9:37:07 will go with 9:37. Since the times are on the index, we pass `left_index` and `right_index`, as we did with merges earlier this chapter:

```
In [42]: pd.merge_asof(
    fb_prices, aapl_prices,
    left_index=True, right_index=True, # datetimes are in the index
    # merge with nearest minute
    direction='nearest', tolerance=pd.Timedelta(30, unit='s')
).head()
```

```
Out[42]:
```

	FB	AAPL
date		
2019-05-20 09:30:00	181.6200	183.5200
2019-05-20 09:31:00	182.6100	NaN
2019-05-20 09:32:00	182.7458	182.8710
2019-05-20 09:33:00	182.9500	182.5000
2019-05-20 09:34:00	183.0600	182.1067

If we don't want to lose the seconds information with the Apple data, we can use `pd.merge_ordered()` instead, which will interleave the two. Note this is an outer join by default (`how` parameter). The only catch here is that we need to reset the index in order to join on it:

```
In [43]: pd.merge_ordered(
    fb_prices.reset_index(), aapl_prices.reset_index()
).set_index('date').head()
```

```
Out[43]:
```

	FB	AAPL
date		
2019-05-20 09:30:00	181.6200	183.520
2019-05-20 09:31:00	182.6100	NaN
2019-05-20 09:31:52	NaN	182.871
2019-05-20 09:32:00	182.7458	NaN
2019-05-20 09:32:36	NaN	182.500

We can pass a `fill_method` to handle NaN values:

```
In [44]: pd.merge_ordered(
    fb_prices.reset_index(), aapl_prices.reset_index(),
```

```
fill_method='ffill')  
.set_index('date').head()
```

Out[44]:

	FB	AAPL
date		
2019-05-20 09:30:00	181.6200	183.520
2019-05-20 09:31:00	182.6100	183.520
2019-05-20 09:31:52	182.6100	182.871
2019-05-20 09:32:00	182.7458	182.871
2019-05-20 09:32:36	182.7458	182.500

Alternatively, we can use fillna() .